

Maksym Yaremchuk

Senior ML Engineer · 7 years · NLP + recsys · PhD candidate

m.yaremchuk@example.com · github.com/myaremchuk · linkedin.com/in/myaremchuk

Summary

ML інженер з 7 років production досвіду. Спеціалізація: NLP (semantic search, RAG, fine-tuning encoders) та recommender systems (two-tower, sequential). Shipped 4 ML systems y production: 2 recsys, 1 semantic search, 1 content classifier. Strong y MLOps — feature stores, model serving, online evaluation.

Experience

Senior ML Engineer — Grammarly (NLP product)

Jul 2023 — Present · Kyiv (remote) · PRODUCTION

- Lead ML engineer on suggestions-quality team: fine-tuned sentence-level encoder для real-time grammar suggestion scoring. Reduced false-positive rate by 23% in A/B test.
- Designed + built feature store on Feast + Snowflake — serves 12 production models, sub-50ms p99 latency.
- Migration of one suggestion model from TensorFlow Serving до Triton Inference Server — 2x throughput, 40% cost reduction on GPU spend.
- Mentor для 3 ML engineers (mid-level). Drive ML system design reviews.

ML Engineer (mid → senior) — Genesis (Universe Group product family)

Feb 2021 — Jun 2023 · Kyiv · PRODUCTION

- Recommender system для consumer subscription app — two-tower retrieval + DLRM ranking on TensorFlow. Lifted retention D7 by 11%.
- Sequential recsys (SASRec-style transformer) для another B2C product — won A/B test against baseline.
- Content classifier (toxicity + spam) — fine-tuned XLM-RoBERTa, deployed via TorchServe.
- Set up Kubeflow Pipelines for offline training, MLflow for experiment tracking.

Junior ML Engineer — Ring Ukraine (Amazon-owned, computer vision)

Aug 2019 — Jan 2021 · Kyiv · PRODUCTION

- Worked on object detection models для smart doorbell devices — YOLOv4 fine-tuning, dataset curation, edge deployment via TensorRT.

Research Intern — Samsung R&D Ukraine

Jun 2019 — Jul 2019 · Kyiv · INTERNSHIP

- Computer vision research — paper reading group, prototype on image segmentation.

Skills

Modeling: Transformers (BERT, RoBERTa, XLM-R), embeddings (sentence-transformers, ColBERT), encoder fine-tuning (contrastive, triplet loss), classification, NER, sequence models, two-tower retrieval, DLRM, SASRec, BERT4Rec.

Frameworks: PyTorch (primary, daily), HuggingFace Transformers, TensorFlow 2.x, JAX (light use), scikit-learn, XGBoost, LightGBM.

MLOps: MLflow, Kubeflow Pipelines, Vertex AI, SageMaker (basic), Triton Inference Server, TorchServe, ONNX, model quantization, Feast feature store.

Languages: Python (advanced, daily 6+ years), SQL (advanced), Bash, basic Go for serving infra.

Data + Infra: Spark (PySpark), Airflow, Databricks, Snowflake, BigQuery, Postgres, Redis (vector + cache), Pinecone (vector DB), AWS (EC2/S3/SageMaker), GCP (Vertex/BigQuery).

Evaluation: offline metrics (NDCG, MRR, AUC, F1), online A/B test design, counterfactual evaluation basics, fairness audits.

Education

NTUU Kyiv Polytechnic Institute (KPI)

PhD in Computer Science (Machine Learning) · 2021 — In Progress (target 2026)

PhD topic: efficient retrieval-augmented generation for low-resource languages. 2 conference papers (EMNLP Findings 2023, COLING 2024). Advisor: Prof. Y.S.

NTUU Kyiv Polytechnic Institute (KPI)

Master's in Applied Math · 2017–2019 · Completed

NTUU Kyiv Polytechnic Institute (KPI)

Bachelor's in Applied Math · 2013–2017 · Completed

Languages

- Ukrainian — native
- English — C1 (conference presentations, paper writing)
- German — A1

Target Roles

Senior / Staff ML Engineer · Applied Scientist · Senior Research Engineer. Open to: NLP, RAG/LLM applications, recsys. Не зацікавлений у pure data-analyst ролях або data-engineer-only.