

Ihor Mazur

Senior Data Engineer · 7 years · Spark / Airflow / Snowflake · Kyiv

i.mazur@example.com · github.com/imazur · linkedin.com/in/imazur

Summary

Data engineer з 7 років production досвіду. Спеціалізація: large-scale batch + streaming pipelines, dimensional modeling, performance tuning warehouse queries. Останні 3 роки — owner data platform у B2B SaaS company (~80 daily users, ~12 TB warehouse). Strong у Spark, Airflow, dbt, Snowflake, Kafka.

Experience

Senior Data Engineer / Data Platform Lead — Restream (B2C streaming SaaS)

Oct 2023 — Present · Kyiv (remote) · PRODUCTION

- Owner data platform: 200+ dbt models, 60+ Airflow DAGs, ingestion з Postgres CDC (Debezium → Kafka → Snowflake).
- Migrated legacy nightly batch (Airflow + Python scripts) до dbt + incremental models — pipeline time 6h → 45min.
- Built streaming consumer service (Kafka → Snowflake via Snowpipe Streaming) for real-time event analytics; lag p99 = 8 seconds.
- Cost optimization: reduced Snowflake monthly spend by \$4.2k (24%) через warehouse-size right-sizing + clustering keys on top 3 fact tables.

Data Engineer (mid → senior) — Grammarly

May 2020 — Sep 2023 · Kyiv · PRODUCTION

- Worked on data ingestion для ML training pipelines — Spark (Scala primarily, some PySpark), Delta Lake on AWS S3, Glue catalog.
- Built feature pipelines для recsys models — pulling from event streams + dimensional dim tables, materialized у Feast.
- Migrated legacy Redshift warehouse до Delta Lake + Trino → cost ↓ 35%, query latency p95 ↓ 60%.

Data Engineer — Genesis (consumer apps)

Aug 2018 — Apr 2020 · Kyiv · PRODUCTION

- Apache Airflow pipelines для marketing analytics (AppsFlyer + Amplitude → BigQuery).
- Built attribution unit-economics model used by growth team — joined cohorts × LTV × CAC across 6 apps.

Junior Data Engineer — SoftServe (outsource)

Jul 2017 — Jul 2018 · Lviv · PRODUCTION

- First commercial role — Python ETL scripts, SQL Server → Hadoop migration for retail client.

Data Engineering Intern — EPAM

Feb 2017 — May 2017 · Lviv · INTERNSHIP

- 3-month bootcamp — Hadoop + Hive + Sqoop fundamentals.

Skills

Languages: Python (advanced, daily 7 years — pandas, PySpark, FastAPI for internal tools), SQL (advanced — window functions, recursive CTEs, query optimization), Scala (intermediate, daily 2 years at Grammarly), Bash, basic Java (Kafka clients).

Warehousing: Snowflake (advanced — clustering, query profiling, resource monitors, Snowpipe + Snowpipe Streaming), BigQuery (mature), Redshift (legacy maintenance), Databricks Delta Lake.

Processing: Spark (Scala + PySpark, batch + structured streaming), dbt (advanced — daily; tests, snapshots, exposures, semantic layer awareness), Trino/Presto, Apache Beam (light).

Orchestration: Airflow (5+ years, Kubernetes executor, sensors, dynamic task mapping), Dagster (evaluated, used on 1 sub-project), Prefect (basic).

Streaming + CDC: Apache Kafka (5+ years, producer/consumer Java + Python clients, Kafka Connect), Debezium for Postgres CDC, AWS Kinesis (limited).

Data modeling: Kimball dimensional (star/snowflake), Data Vault basics, slowly changing dimensions (Type 1/2/3), incremental materializations in dbt.

Cloud: AWS (S3, Glue, EMR, Lambda, IAM, EventBridge), GCP (BigQuery, Cloud Storage, Pub/Sub).

Tools: Git, GitHub Actions, Terraform (basic — DAG infra), Looker (semantic layer), Mode Analytics, dbt-power-user VS Code extension, Great Expectations + Monte Carlo (data quality).

Governance + obs: dbt tests, Great Expectations, Monte Carlo Data Observability, lineage via dbt + OpenLineage.

Education

NTUU Kyiv Polytechnic Institute (KPI)

Master's in Computer Science · 2015-2017 · Completed

NTUU Kyiv Polytechnic Institute (KPI)

Bachelor's in Computer Science · 2011-2015 · Completed

Certifications: SnowPro Core (2023), AWS Data Engineer — Associate (2024), dbt Cloud Analytics Engineering Certificate (2022).

Languages

- Ukrainian — native
- English — C1 (daily — written design docs + remote standups)

Target Roles

Senior Data Engineer · Staff Data Engineer · Data Platform Engineer · Analytics Engineering Lead. Цікавлять product companies з real-time data needs. Не зацікавлений у Data Analyst ролях (BI dashboards only).